

# **COP4635: Systems & Networks II**

Network Programming:  
Sockets & Client Server


- Sockets
  - communication endpoints for processes
  - can be on different machines
  - can be used over network
- Client - Server
  - client is a process sending requests to a server
  - server is a continuously running process responding to client requests
  - client – server use connection-oriented service

- Socket
  - Communication endpoint
  - Commonly used for network communication
  - Composed of three parts
    - **IP Address**: Identifies machine
    - **Port**: Identifies process on machine
    - **Protocol**: Specifies communication parameters, order, etc.

# Network Communication

Machine 1

Machine 2



The following data types are needed to build sockets:

<code>u_char</code>	unsigned char (8-bit)
<code>u_short</code>	unsigned short (16-bit)
<code>u_long</code>	unsigned long (32-bit)

- Specifies machine on Internet
  - Assume we are using IPv4 (not IPv6)
  - 32-bit integer number
  - Usually written as dotted quad
  - Type is `u_long` (unsigned long)
  - Wrapped in structure

```
struct in_addr
{
    u_long    s_addr;
};
```

# Port Numbers

Logical not physical  
16-bit integer  
(short)

0-1023	Well-known (reserved)
1024-49151	Registered (discouraged)
49152-65535	Dynamic/Private

7	echo		25	smtp		113	ident
11	systat		37	time		119	nntp
21	ftp		79	finger		135	RPC
22	ssh		80	http		143	imap
23	telnet		110	pop3		443	https

- Specifies
  - Format of messages
  - Headers
  - Other communication parameters
  - A few examples:

PF\_INET

PF\_INET6

PF\_BLUETOOTH

PF\_APPLETALK



```
struct sockaddr_in
{
    u_char                sin_len;
                        // normally
    not used
    u_short               sin_family;
    u_short               sin_port;
    struct in_addr        sin_addr;
    char                 sin_zero[8];
                        // normally
    not used
};
```

## TCP

- Connection-oriented, reliable
- Also called *stream*

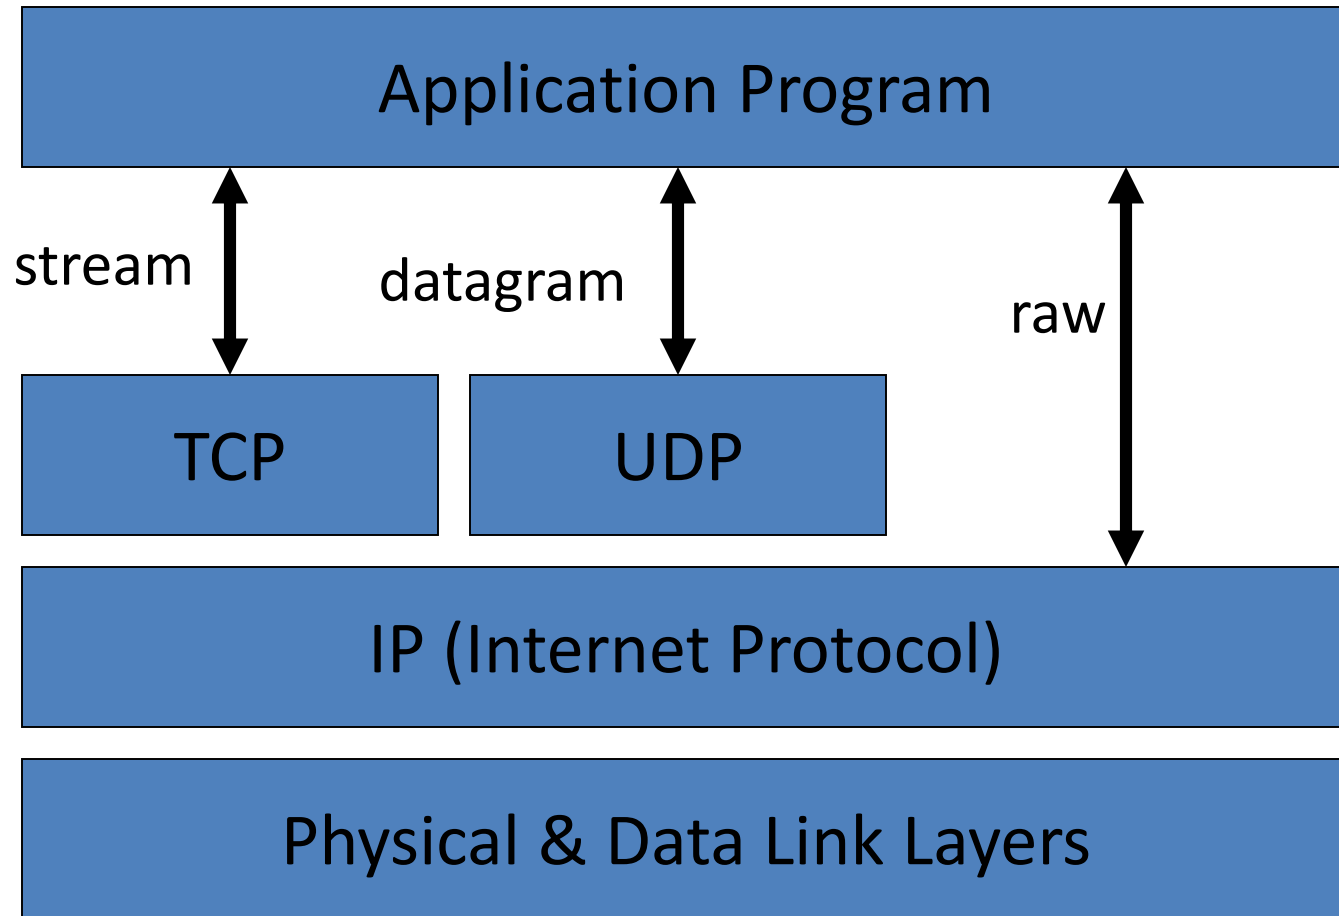
## UDP

- Connectionless, unreliable
- Also called *datagram*

## RAW

- Bypasses niceties of above
- Functions at IP level

# Protocol Layers



All machines are not created equal.

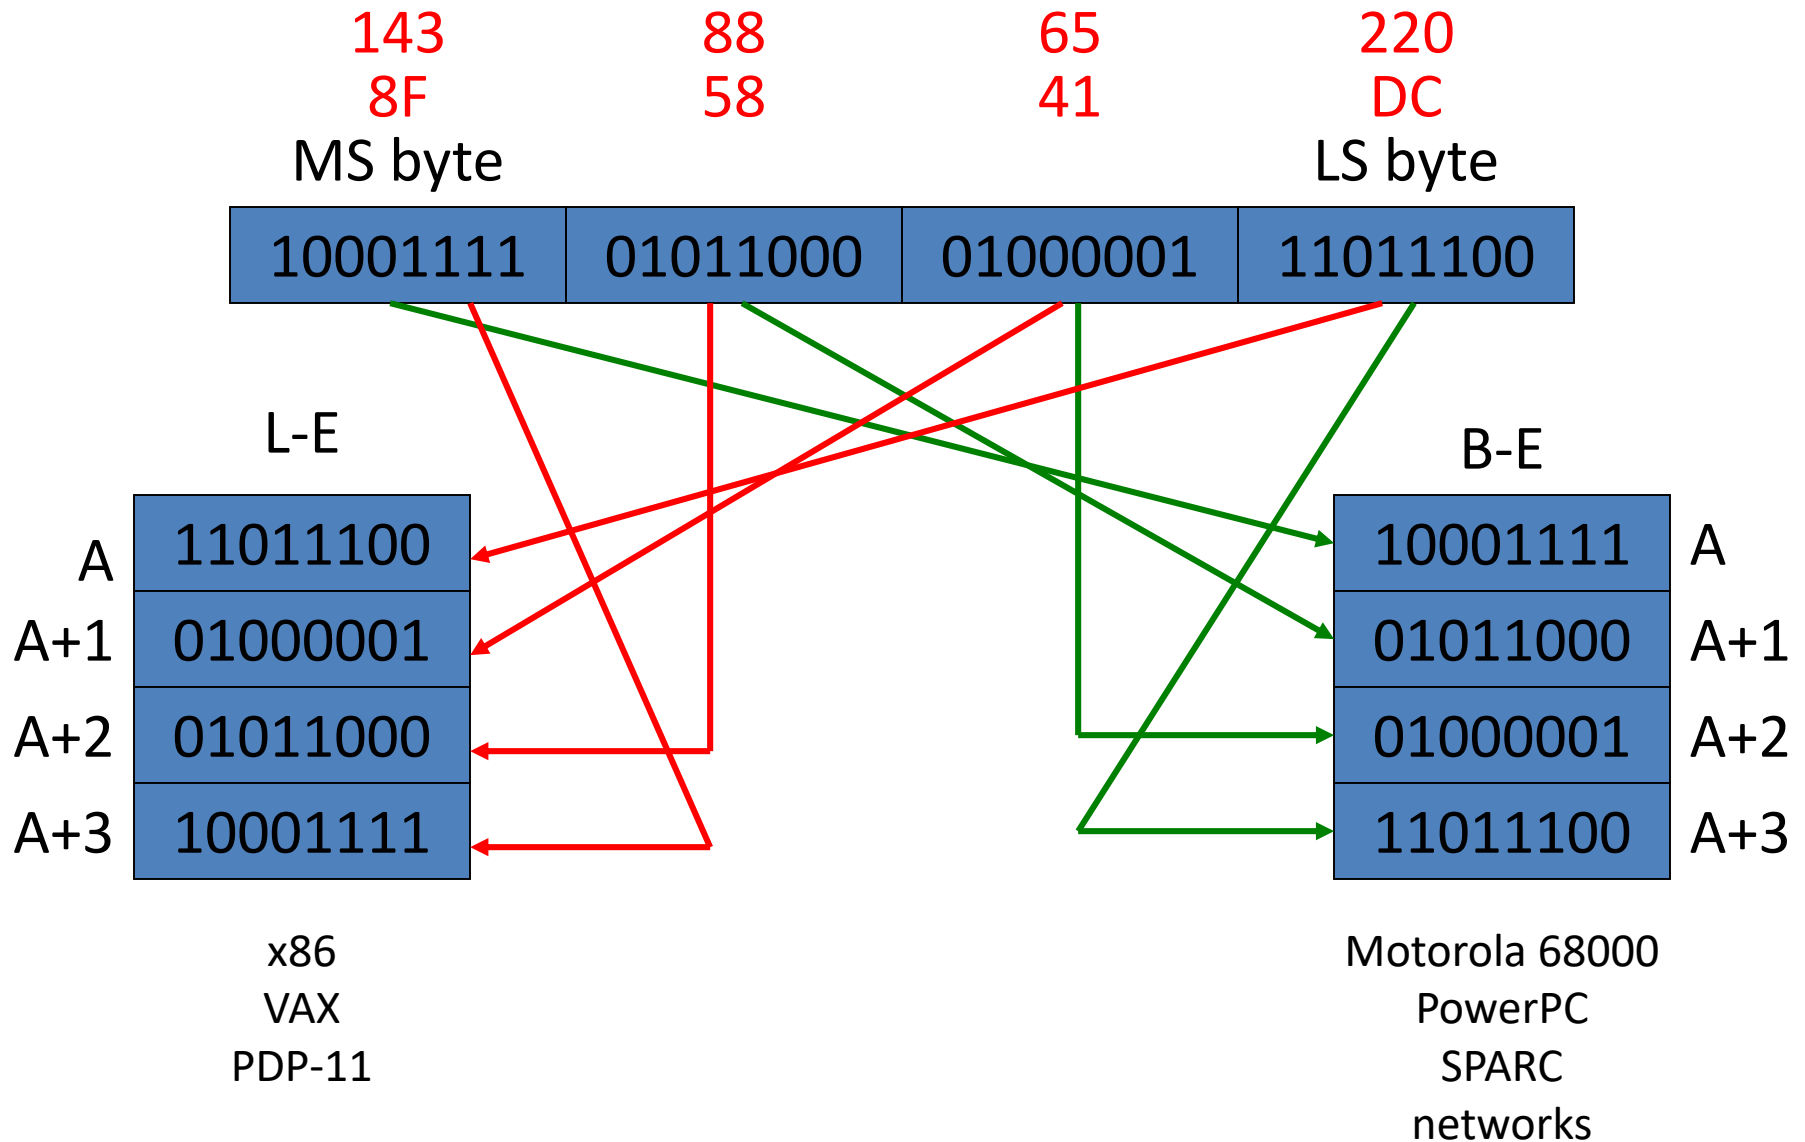
2- & 4-byte values can be stored differently.



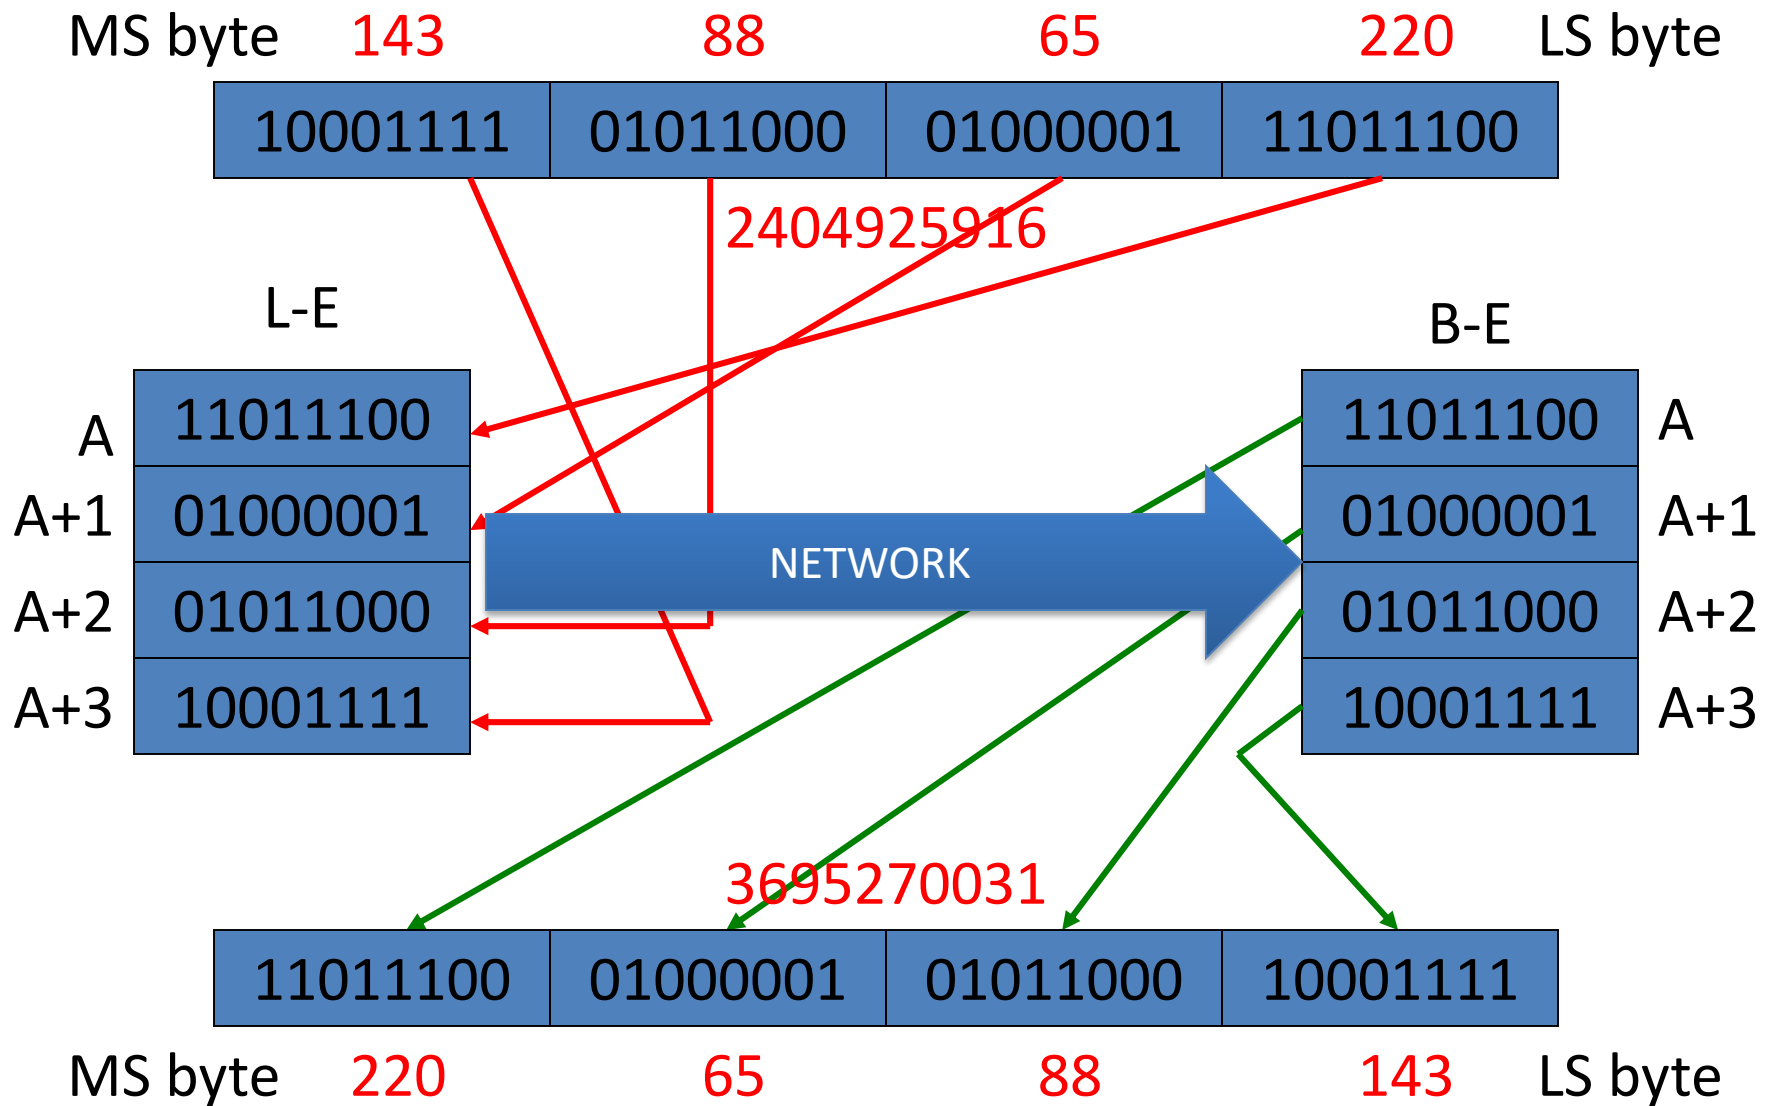
**Endianness** (byte order) can differ.

The very term *big-endian* comes from Jonathan Swift's satiric novel *Gulliver's Travels*, where tensions are described in Lilliput and Blefuscu: whereas royal edict in Lilliput requires cracking open one's soft-boiled egg at the small end, inhabitants of the rival kingdom of Blefuscu crack theirs at the big end (hence the name *Big-endians*).

# Byte Ordering



# The Problem



## Machine A

- Big Endian
- Send 2-byte decimal value 821
- In binary, that's (00000011 00110101)

## Machine B

- Little Endian
- Receive binary (00110101 00000011)
- In decimal, that's 13571

# Network Byte Order

- Network bytes order is big-endian
- Special functions to convert to/from NBO

`htons`    host to network short

`ntohs`    network to host short

`htonl`    host to network long

`ntohl`    network to host long

- Trick
  - B-E machine: functions do nothing
  - L-E machine: functions reverse bytes
  - Always use on data
  - Don't worry about endianness of machine



# Address Transformation

Lots of functions to manipulate IP Addresses.

Refer to man pages for details and include files.

```
int inet_aton( char* cp,  
              struct in_addr *inp)
```

Converts dotted-quad (*cp*) to 32-bit address in network byte order (*inp*)

```
char *inet_ntoa( struct in_addr in)
```

Converts 32-bit address in network byte order (*in*) to dotted-quad (return pointer)

# Host Information

Given hostname, we need to lookup the IP address:

```
struct hostent * gethostbyname( char* name)

struct hostent
{
    char      *h_name;           /* official name */
    char      **h_aliases;       /* list of aliases */
    int        h_addrtype;       /* PF_INET */
    int        h_length;         /* length, e.g. 4 */
    char      **h_addr_list;     /* list of addr. */
};
```

# Creating a Socket

Access similar to pipe or file

Through file descriptor table

Called *socket descriptor* in this case

```
int socket(  
    int domain,    /* PF_INET */  
    int type,      /* SOCK_DGRAM, SOCK_STREAM */  
    int protocol /* 0 */  
);
```

Return value of  $-1$  indicates an error

# Putting It All Together

1. Create a socket
2. Get information about *dest* host.
3. Fill in *dest addr* structure.
4. Connect to remote host.
5. Send /receive data to *dest*.
6. Close connection.

# Create a socket

```
socketFD = socket( AF_INET, SOCK_DGRAM, 0 );  
if ( socketFD < 0 ) {  
    perror( "sendUDP:socket" );  
    return -1;  
}
```

Obtaining information about a remote host:

```
hostptr = gethostbyname( argv[1] );  
if ( hostptr == NULL ) {  
    perror( "sendUDP:gethostbyname" );  
    return -1;  
}
```

Function makes a call to DNS and returns a struct called *hostent*.

# Fill In Dest Addr

```
memset( &destaddr, 0,  
        sizeof(struct sockaddr_in) );  
destaddr.sin_family = AF_INET;  
memcpy( &destaddr.sin_addr,  
        hostptr->h_addr,  
        hostptr->h_length );  
destaddr.sin_port =  
    htons( (u_short)atoi(argv[2]) );
```

# Send Message

```
fgets( buffer, 256, stdin );  
sendlen = strlen( buffer ) - 1;  
bytes = sendto(  
    socketFD,  
    buffer,  
    (size_t)sendlen,  
    0,  
    (struct sockaddr *)&destaddr,  
    (socklen_t)sizeof(destaddr)  
);
```



# Receiver Extra Step

- Receiver must have port number.
- Delivery:
  - Packet → Machine (host) → OS → Process
- The `bind()` system call provides “name” for socket.
- Can send without a “named socket”, but must have “named socket” to receive.
  - Similar to U.S. Postal Service

```
int bind(  
    int sockfd,  
    struct sockaddr *my_addr,  
    socklen_t addrlen  
);
```

# What Is sockaddr ?

```
struct sockaddr_in
{
    u_char            sin_len;
    u_short           sin_family;
    u_short           sin_port;
    struct in_addr    sin_addr;
    char              sin_zero[8];
};
```

`sendto()` → `recvfrom()`

Sender needs to send to a socket

- IP Address
- Protocol Family
- Port number

Receiver needs to establish socket

- Needs to establish port number
- How?

If 1<sup>st</sup> action is receive,

- Must establish port number

- And announce port number

Why?

- So sender knows where to send

Otherwise

- Receiver is “*out there*” somewhere

## Static binding

- **User** provides port number; some options
- Hardcoded in program

```
DA.sin_port = htons ( 51664 );
```

- Provided on command-line

```
DA.sin_port =  
    htons ( (u_short) atoi (argv[1]) );
```

- Provided via user input

```
DA.sin_port =  
    htons ( (u_short) atoi (inString) );
```

## Dynamic binding

- **Kernel** provides port number
- Send special “you pick” to kernel
- Value of 0 (zero) is indicator

```
DA.sin_port = 0;
```

- Kernel picks next available port number
- Stores port number internally

# Example

```
int bind(  
    int sockfd,  
    struct sockaddr *my_addr,  
    socklen_t addrlen  
);
```

The `bind()`  
selects next  
available port

```
int getsockname(  
    int sockfd,  
    struct sockaddr *my_addr,  
    socklen_t *addrlen  
);
```

The  
`getsockname()`  
returns the *name* of  
the socket



# Announcing Socket Name

```
printf("name: %s\n",  
      hostptr->h_name);
```

```
printf("addr: %s\n",  
      inet_ntoa(srcAddr.sin_addr));
```

```
printf("port: %d\n",  
      ntohs(srcAddr.sin_port));
```

```
> ./recvUDP
```

```
name: cs-ssh1.cs.uwf.edu
```

```
addr: 143.88.64.151
```

```
port: 54464
```

```
> ./recvUDP
```

```
name: cs-ssh1.cs.uwf.edu
```

```
addr: 143.88.64.151
```

```
port: 54208
```

```
>
```

Do you need a port number to send?

*Yes*

Where do you get it?

*Dynamically – from kernel*

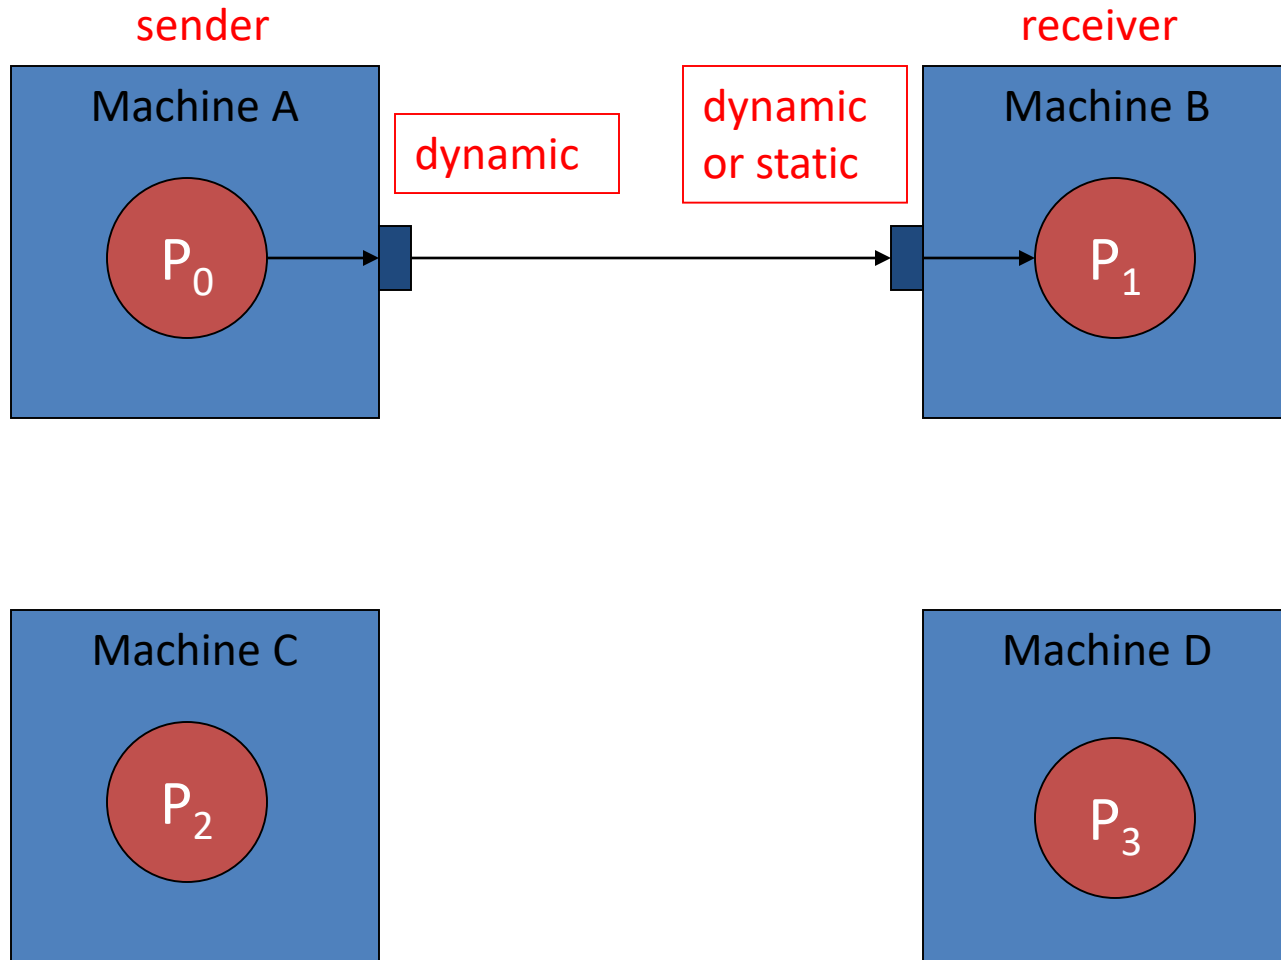
How?

`sendto()` allocates port if needed

This is second form of dynamic binding

Try using `sendto()` then `getsockname()`

# Big Picture - A



Sit at **ONE** machine (machine A)

Open a terminal

- Terminal displays on A
- Commands run on A

Open a second terminal

- Terminal displays on A
- Commands run on A
- Use `ssh` to connect to machine B
- Terminal displays on A
- Commands run on B

# What About My Files?

## Multiplatform Lab & SSH servers use NFS

- Network File System
- Files stored on “file server”
- Files available on all machines
- Save on A, new file also visible on B

Compile on A, execute on B

display

**sse-250-061**

```
> ./recvUDP
name: sse-250-061.cs.uwf.edu
addr: 143.88.64.206
port: 13925
got [hello]
```

execute

**sse-250-061**

**sse-250-061**

**sse-250-05e**

```
> ssh sse-250-05e.cs.uwf.edu
> ./sendUDP sse-250-061.cs.uwf.edu 13925
Send this> hello
sent 5 bytes
>
```

## Requires endpoints (sockets)

- Protocol (AF\_INET)
- IP Address (32-bit or dotted quad)
- Port Number (16-bit; static or dynamic)

## UDP

- Sends *datagrams*
- May arrive; may be lost
- May arrive in different order than sent



TCP is a protocol designed for the reliable transfer of bytes from one process to another over a network

How?

- Sequence numbers
- ACK
- Timers (what if small file?)

Designed for client/server model

Server is at well-known place

Usually only one server

Clients (multiple) are anywhere

Client contacts server for *service*

Client and server communicate to facilitate service

Client and server disconnect

Server continues to wait for next client contact

Client and Server have “1-to-1” connection

No other process can use connection

Connection is bidirectional

- Client can send to server
- Server can send to client

Requires buffers, segments, handshakes, ...

All (sender, receiver, client, server)

- `socket()`
- `gethostbyname( who )`

TCP must create client/server connection

Client

`connect()`

Server

`bind()`

`listen()`

`accept()`

`bind()`

- Makes receipt before send possible

`listen()`

- Expresses willingness to accept incoming connections
- Sets connection parameters (queue limit, etc.) for incoming connections


`accept()`

- Accepts a connection on a listening socket
- Returns file descriptor for newly created 1-to-1 socket


# TCP Handshake

- Done by `connect ()` / `accept ()`
- Client
  - Sends “CONN REQ” to server socket
- Server
  - Accepts request
  - Creates new socket (dynamic port binding)
  - Randomly generates 1st  $S \rightarrow C$  sequence number
  - Sends new socket & 1st  $S \rightarrow C$  sequence number to client (“CONN ACK”)
- Client
  - Accepts message
  - Randomly generates 1st  $C \rightarrow S$  sequence number
  - Sends 1st  $C \rightarrow S$  sequence number to server (“CONNECT”)

# Client to Server




# One-to-One Communication






# Next Client



# Still One-to-One



`socket()`

`gethostbyname( server )`

`connect()`

`send() / recv()`

`close()`

`socket()`

`gethostbyname( self )`

`listen()`

`accept()`

`send() / recv()`

`close()`


# Normal Server Operation

```
listen()  
while ( 1 ) {  
    newFD = accept()  
    pthread_create()  
}
```

## Child thread

- Handles request
- Communicates with client
- Terminates when service complete
- May block for I/O

# Threaded Server



## Sender

```
socket ()
gethostbyname ( dest )
sendto ()
close ()
```

## Receiver

```
socket ()
gethostbyname ( dest )
bind ()
recvfrom ()
close ()
```



Datagram is limited in size (i.e., 10K)

How can we transfer large file (i.e., 100K)?

## UDP Sender

- Partition file into 10 datagrams
- Send datagram 1, then datagram 2, ...

## UDP Receiver

- Receive all datagrams
- Reconstruct file from datagrams



Some datagrams could be lost

1 2 5 6 8 9 10

Datagrams could arrive out of order

1 2 5 8 10 3 9 7 4 6

Or Both

1 5 8 10 7 4 6

What now?

- We didn't receive file as it was sent

Add a number to each datagram

- Sequence number (start at 0)
- Receiver can reorder datagrams
- Receiver can ask for lost datagram to be resent
- Receiver can send acknowledgement (ACK) of receipt for datagrams

# Creating Streaming Sockets in Java

- A socket is created using the `java.net.Socket`:

```
Socket socket = new Socket(host, port);
```

`host` is a string specifying a host; can be `localhost`

`port` is the port number on which the host is contacted

- **Example:**

```
Socket socket = new  
    java.net.Socket(localhost, 8080);
```

# Creating Datagram Sockets in Java

- A datagram socket is created using `java.net.DatagramSocket`:  

```
DatagramSocket dsocket =  
    new DatagramSocket(port);
```

`port` is the UDP port on which packets will be received;  
the datagram socket is bound to host's Internet address
- Example:  

```
DatagramSocket dsocket =  
    java.net.DatagramSocket(8080);
```

- IO from a socket stream is like file IO:

```
// the buffer that stores read data
byte[] buffer = new byte[1024];

// creating an input stream; an error occurs if the
// socket is not open or not connected
InputStream iStream = socket.getInputStream();
OutputStream oStream = socket.getOutputStream();

// sending data stored in the buffer
oStream.write(buffer)

// reading data stored in the buffer
int numBytesRead = iStream.read(buffer)
```

- To end a connection, either side of the end can call `close()`.

For example:

```
socket.close();
```

- If one end closes the socket and the other end does a
  - `read()`, the function will throw an IO exception,
  - `write()`, the function will throw an IO exception.

- Server sockets in Java

(`java.net.ServerSocket`):

- A server socket is created as follows:

```
ServerSocket ssocket = new  
    ServerSocket(port);
```

`port` is the port on which the server listens for connections

- Example:

```
ServerSocket ssocket = ServerSocket(8080);
```

# Designing a Network Protocol

- Server and client respond to messages based on a well-defined protocol.
- Protocol defines actions and data that a client requests from a server.
- Client and server compile and parse message based on protocol.
- Message can be in different formats including text, binary, or mixed.



# Network Protocol Examples

- Online banking that supports users perform various banking transactions:  
*<action>operation</action><value>aValue</value>*
- Browser sends message to a server to retrieve a Web page.

Client sending a message to the server:

```
private String sendMessage(String message) throws Exception {
    Socket socket;
    InputStream iStream;
    OutputStream oStream;
    try {
        // creates socket to communicate with server
        socket = new Socket(_host, _port);

        // access data streams from socket
        iStream = socket.getInputStream();
        oStream = socket.getOutputStream();
    }
    catch (Exception xcp) {
        throw new Exception ("connection error: unable to connect to the server");
    }
    .....
}
```

## Client sending a message to the server:

```
String response;
try {
    PrintWriter printWriter = new PrintWriter(oStream);
    Scanner scanner = new Scanner(iStream);

    // sending message
    printWriter.println(message);
    printWriter.flush();

    // receiving response
    response = scanner.nextLine();
}
catch (Exception xcp) {
    socket.close();
    throw new Exception ("unable to communicate with server");
}
// close connection to server
socket.close();

return response;
}
```

PrintWriter and Scanner  
aid in writing and reading  
text from a socket.

## Sample protocol for banking transaction:

Client sends a message to the server to create an account:

`create | name | deposit`

*name*: name of the account owner

*deposit*: initial deposit of the account

Example:

`create | joe smith | 300`

Server sends message confirming account creation:

*account\_number*

*account\_number*: the number of the account

Example:

5

## Server socket waiting for incoming connection:

```

ServerSocket serverSocket;
try {
    serverSocket = new ServerSocket(_port);
}
catch (Exception xcp) {
    throw new Exception ("unable to start server on port " + _port);
}
while (true) {
    Socket socket = serverSocket.accept();
    try {
        // get communication channels of socket
        InputStream iStream = socket.getInputStream();
        OutputStream oStream = socket.getOutputStream();

        // start a request handler in a separate thread
        RequestHandler requestHandler = new RequestHandler(iStream, oStream);
        new Thread(requestHandler).start();
    }
    catch (Exception xcp) {
        // do nothing; client has dropped connection
    }
}

```

Socket is accepting a connection to a host.

Handling request from client in a separate thread.

Handling a client's request:

```
public void run() {  
    // read data from  
    String requestMessage = _scanner.nextLine();  
    String response = handleRequest(requestMessage);  
  
    // send response back to host  
    sendMessage(response);  
  
    // close communication channels  
    _printWriter.close();  
    _scanner.close();  
}
```

Read request from client and handle it.

Sends computed response.

Closes communication.

Creating a bank account using information submitted by client:

```
public int createAccount(String name, float initialBalance) throws Exception {  
    int accountNumber = _bankData.getNextAccountNumber();  
    BankAccount bankAccount = new BankAccount(accountNumber, name, initialBalance);  
    _bankData.write(bankAccount);  
  
    return accountNumber;  
}
```

Sends message back to client:

```
public void sendMessage(String message) {  
    try {  
        _printWriter.println(message);  
        _printWriter.flush();  
    }  
    catch (Exception xcp) {  
        System.err.println("error occurred sending message to host: " + message);  
    }  
}
```

- Server accepts incoming connections continuously.
  - `accept()` blocks server until a new connection is requested
  - requests from clients are handled in separate threads
- Server terminates when program is shutdown.
  - more graceful shutdown could be done through a special server administration console



# Summary Questions

- What is different between a UDP and a TCP socket?
- What values have to be mapped between host and network to maintain a correct byte order?
- When is *bind()* needed, when is it not needed to communicate with a host?
- What is the purpose of calling *listen()* on a TCP server socket?
- If a TCP server program communicates with 3 TCP clients how many sockets does the server program have open?